

# Classification on Immunotherapy data set

Munkhsukh Munkhtsetseg

Eötvös Loránd University, H-1053 Budapest, Egyetem tér 1-3  
<https://www.elte.hu/>

**Abstract.** The goal of this paper is to contribute to a clear understanding of a classification task on Immunotherapy data set. We argue that the classification task can be considered different machine learning algorithms and techniques to improve our result of ML algorithms.

**Keywords:** Machine learning · Immunotherapy data set · Classification.

## 1 Introduction to Data set

### 1.1 Description

Immunotherapy dataset [1] comes from the UCI Machine Learning repository, and the dataset contains information about wart treatment results of 90 patients using immunotherapy. For number of records, It is too low number to train ML model though, We will examine to create a back-bone.

**Table 1.** Immunotherapy data set description table

|                            |                |                       |            |
|----------------------------|----------------|-----------------------|------------|
| Data Set Characteristics:  | Univariate     | Number of Instances:  | 90         |
| Attribute Characteristics: | Integer, Real  | Number of Attributes: | 8          |
| Associated Tasks:          | Classification | Missing Values?       | N/A        |
| Area:                      | Life           | Date Donated          | 2018-01-04 |


**Table 2.** Immunotherapy data table description

|                     | count | mean      | std        | min  | 25%   | 50%   | 75%     | max   |
|---------------------|-------|-----------|------------|------|-------|-------|---------|-------|
| sex                 | 90.0  | 1.544444  | 0.500811   | 1.0  | 1.00  | 2.00  | 2.0000  | 2.0   |
| age                 | 90.0  | 31.044444 | 12.235435  | 15.0 | 20.25 | 28.50 | 41.7500 | 56.0  |
| Time                | 90.0  | 7.230556  | 3.098166   | 1.0  | 5.00  | 7.75  | 9.9375  | 12.0  |
| Number_of_Warts     | 90.0  | 6.144444  | 4.212238   | 1.0  | 2.00  | 6.00  | 8.7500  | 19.0  |
| Type                | 90.0  | 1.711111  | 0.824409   | 1.0  | 1.00  | 1.00  | 2.0000  | 3.0   |
| Area                | 90.0  | 95.700000 | 136.614643 | 6.0  | 35.50 | 53.00 | 80.7500 | 900.0 |
| induration_diameter | 90.0  | 14.333333 | 17.217707  | 2.0  | 5.00  | 7.00  | 9.0000  | 70.0  |
| Result_of_Treatment | 90.0  | 0.788889  | 0.410383   | 0.0  | 1.00  | 1.00  | 1.0000  | 1.0   |

It includes 8 fields and sex, type and result of treatment fields are categorical. Other age, time, number of warts, area and induration diameter fields are numerical. In sex field, 1 and 2 represents male and female, and in result of treatment field, 1 means treatment succeeded, and 0 means opposite. Last one, there are 3 types of treatment.

## 1.2 Data exploration

This section we will explore the results.



**Fig. 1.** In our data set, 71 treatment has passed, and 19 ones are failed.

The ratio is 71:19, which means our classes are significantly unbalanced.

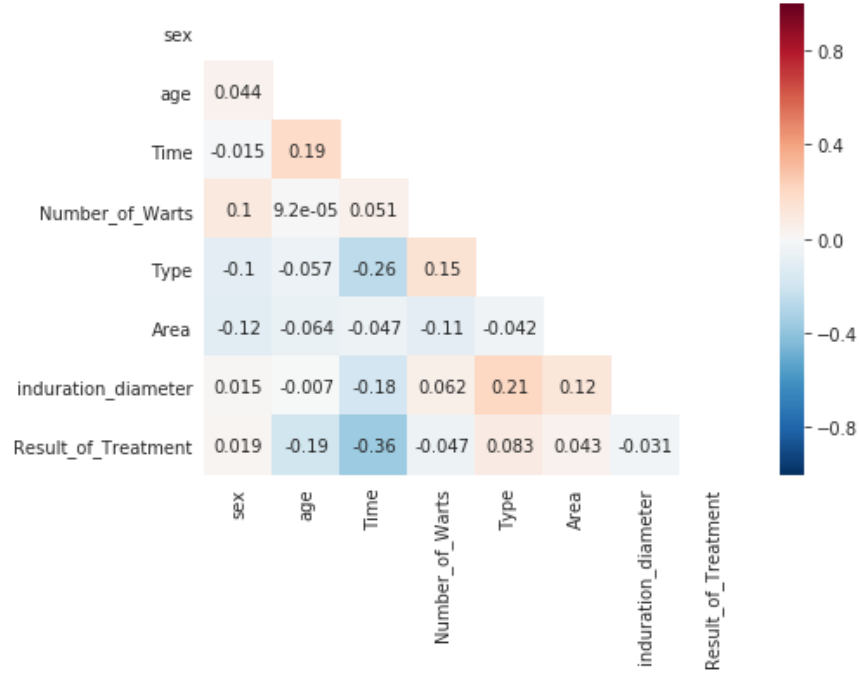
**Table 3.** Categorical means by result of treatment and difference between them.

|                     | 0         | 1         | Difference |
|---------------------|-----------|-----------|------------|
| sex                 | 1.526316  | 1.549296  | 0.022980   |
| age                 | 35.473684 | 29.859155 | -5.614529  |
| Time                | 9.381579  | 6.654930  | -2.726649  |
| Number_of_Warts     | 6.526316  | 6.042254  | -0.484062  |
| Type                | 1.578947  | 1.746479  | 0.167532   |
| Area                | 84.315789 | 98.746479 | 14.430689  |
| induration_diameter | 15.368421 | 14.056338 | -1.312083  |

*Observation* Table 3 gives following observations. Younger people tend to pass treatment, while less time and bigger area can give us better results.

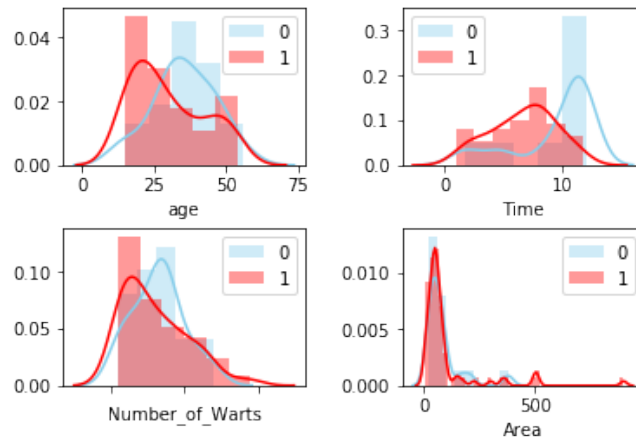
### 1.3 Visualization

In order to explore more, lets see the distribution of data and the correlations.




**Fig. 2.** Correlation diagram of data table

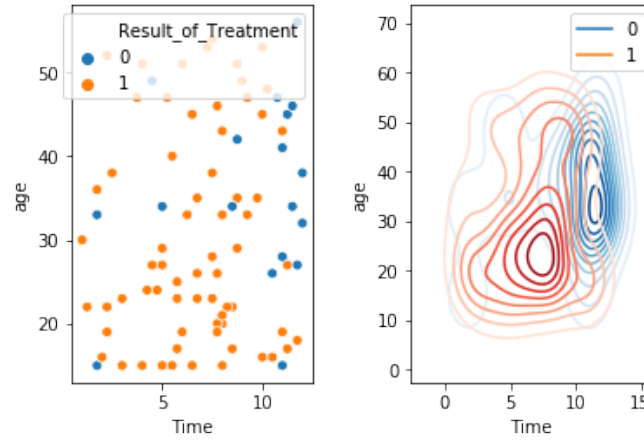
*Observation* Data columns with very similar trends are also likely to carry very similar information. Here we calculate the correlation coefficient between numerical and nominal columns as the Coefficient. Here we can Time field has high negative correlations with result and type fields. We may use this diagram to reduce features later.




**Fig. 3.** Histogram diagrams in age, time, number of warts and area by the result of treatment. Age, time and number of warts might be strong predictors.



**Fig. 4.** Histogram diagrams in age, time and their relationship. As we see, around 7.5 and 20 is the most frequent time and age.



**Fig. 5.** Here we can see their relationship on the result of treatment. Age and time fields can be strongest predictors.



**Fig. 6.** As the correlation diagram, Time and type had a high score. Here, we can see their relation.

## 2 Model

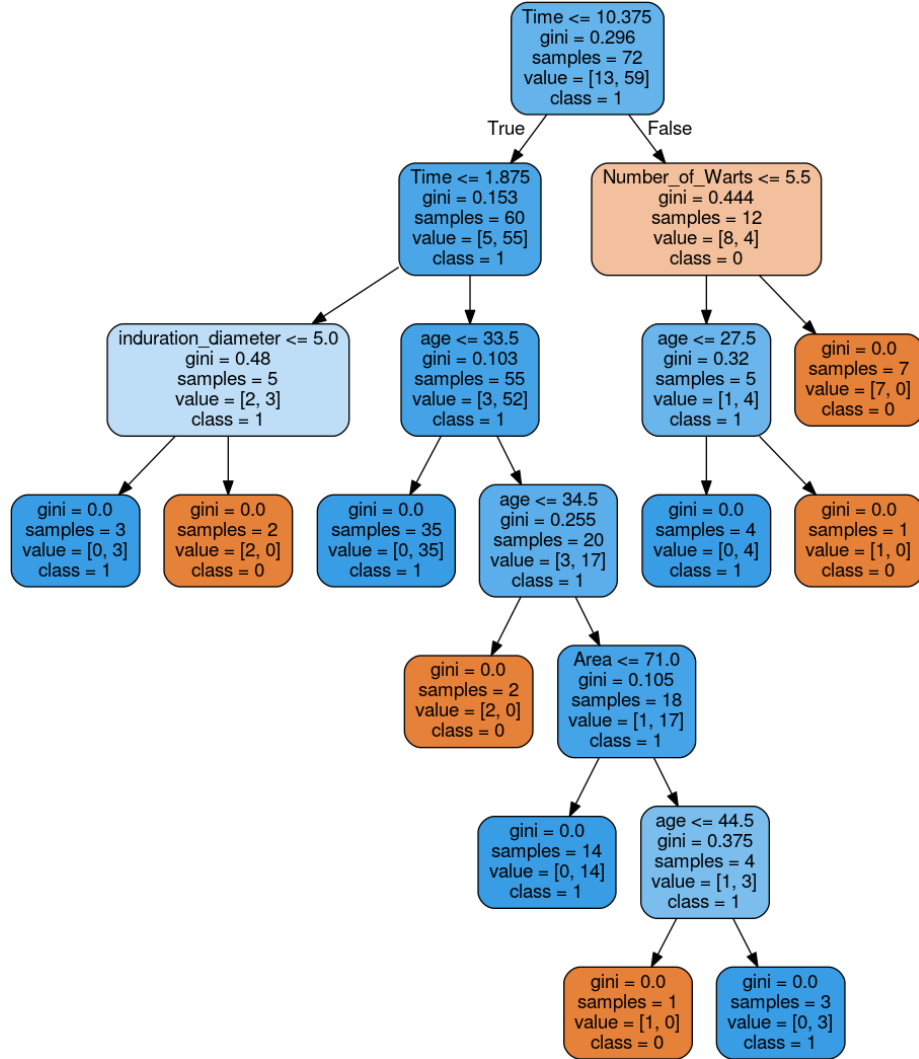
*Train/Test Split* First of all, the data we use is usually split into training data and test data. The training set contains a known output and the model learns on this data in order to be generalized to other data later on. We have the test dataset (or subset) in order to test our model's prediction on this subset. Here, we split our samples into 20/80 ratio test and training samples, respectively.

### 2.1 Decision tree

Given a data of attributes together with its classes, a decision tree produces a sequence of rules that can be used to classify the data.

*Description:* Decision Tree [2], as its name says, makes decision with tree-like model. It splits the sample into two or more homogeneous sets (leaves) based on the most significant differentiators in your input variables. To choose a differentiator (predictor), the algorithm considers all features and does a binary split on them (for categorical data, split by cat; for continuous, pick a cut-off threshold). It will then choose the one with the least cost (i.e. highest accuracy), and repeats recursively, until it successfully splits the data in all leaves (or reaches the maximum depth).

Another reason I chose decision tree model is easy to visualize and examine model. Following figure shows us the result.




**Fig. 7.** Decision tree model after trained with our training data.

## 2.2 Logistic regression

Logistic regression is a classification algorithm used to assign observations to a discrete set of classes. We will use binary logistic regression. In order to map predicted values to probabilities, we use the sigmoid function. The function maps any real value into another value between 0 and 1. In machine learning, we use sigmoid to map predictions to probabilities.

$$S(z) = \frac{1}{1 + e^{-z}} \quad (1)$$



**Fig. 8.** Sigmoid function

Apply Sigmoid function on linear regression:


$$S(z) = \frac{1}{1 + e^{-b_0 + b_1 x_1 + b_2 x_2 + \dots + b_n x_n}} \quad (2)$$

If the weighted sum of inputs is greater than zero, the predicted class is 1 and vice-versa. So the decision boundary separating both the classes can be found by setting the weighted sum of inputs to 0.

**Cost function** <https://www.overleaf.com/project/5e308d9f1394c60001211bc5>  
The cost function for a single training example can be given by:

$$cost = \begin{cases} -\log(h(x)) & \text{if } y = 1, \\ -\log(1 - h(x)) & \text{if } y = 0 \end{cases}$$





**Fig. 9.** Cost function

If the actual class is 1 and the model predicts 0, we should highly penalize it and vice-versa. We can combine both of the equations using:

$$cost(h(x), y) = -y \log(h(x)) - (1 - y) \log(1 - h(x))$$

The cost for all the training examples denoted by  $J()$  can be computed by taking the average over the cost of all the training samples


$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^i \log(h(x^i)) + (1 - y^i) \log(1 - h(x^i))]$$

where  $m$  is the number training samples. We will use gradient descent to minimize the cost function.

Optimization terminated successfully.  
Current function value: 0.473337  
Iterations 6

#### Results: Logit

```
=====
Model:                Logit                Pseudo R-squared: -0.002
Dependent Variable:    Result_of_Treatment  AIC:                82.1605
Date:                  2020-01-30 13:18      BIC:                98.0972
No. Observations:      72                  Log-Likelihood:     -34.080
Df Model:               6                   LL-Null:            -34.001
Df Residuals:           65                  LLR p-value:        1.0000
Converged:              1.0000              Scale:              1.0000
```



**Fig. 10.** Sigmoid function on on our training samples

No. Iterations: 6.0000


|                     | Coef.   | Std.Err. | z       | P> z   | [0.025 0.975] |        |
|---------------------|---------|----------|---------|--------|---------------|--------|
| sex                 | 0.9095  | 0.5196   | 1.7504  | 0.0800 | -0.1089       | 1.9279 |
| age                 | 0.0210  | 0.0270   | 0.7771  | 0.4371 | -0.0320       | 0.0740 |
| Time                | -0.1318 | 0.0956   | -1.3794 | 0.1678 | -0.3192       | 0.0555 |
| Number_of_Warts     | -0.0218 | 0.0741   | -0.2938 | 0.7689 | -0.1669       | 0.1234 |
| Type                | 0.2155  | 0.3432   | 0.6279  | 0.5300 | -0.4572       | 0.8882 |
| Area                | 0.0007  | 0.0024   | 0.2703  | 0.7869 | -0.0041       | 0.0054 |
| induration_diameter | 0.0023  | 0.0213   | 0.1078  | 0.9142 | -0.0394       | 0.0440 |

Now we tried our own code to predict logistic regression with Immunotherapy data set. Later we will use scikit-learn library and compare other algorithms.

```

iter: 0 cost: 1.8732674100249116
iter: 100 cost: 0.6522403704435823
iter: 200 cost: 0.609406517463262
iter: 300 cost: 0.5777195777911466
iter: 400 cost: 0.5551677781320082
iter: 500 cost: 0.5395701581670956
iter: 600 cost: 0.5289746327726486
iter: 700 cost: 0.5218382544817948
iter: 800 cost: 0.517035287249884
iter: 900 cost: 0.5137841005838157

```



**Fig. 11.** Training process

### 2.3 Evaluation of Model

**Confusion matrix** A confusion matrix is a table that is often used to describe the performance of a classification model on a set of test data for which the true values are known.

- true positives (TP): These are cases in which we predicted yes and are actually yes.
- true negatives (TN): We predicted no, and no in actual.
- false positives (FP): We predicted yes, but actual is no. (Type I error)
- false negatives (FN): We predicted no, yes in actual. (Type II error)

| Actual   | Predicted           |                     |
|----------|---------------------|---------------------|
|          |                     |                     |
|          | Positive            | Negative            |
| Positive | True Positive (TP)  | False Negative (FN) |
| Negative | False Positive (FP) | True Negative (TN)  |

**Fig. 12.** Confusion matrix

- Accuracy: Overall, how often is the classifier correct?

$$Accuracy = (TP + TN) / total$$

- Misclassification Rate(Error Rate): Overall, how often is it wrong?

$$MisclassificationRate = (FP + FN)/total$$

- True Positive Rate(Sensitivity or Recall): When it's actually yes, how often does it predict yes?

$$TruePositiveRate = TP/(actual_{yes})$$

- False Positive Rate: When it's actually no, how often does it predict yes?

$$FalsePositiveRate = FP/(actual_{no})$$

- True Negative Rate(Specificity): When it's actually no, how often does it predict no?

$$TrueNegativeRate = TN/actual_{no}$$

- Precision: When it predicts yes, how often is it correct?

$$Precision = TP/predictedyes$$

- Prevalence: How often does the yes condition actually occur in our sample?

$$Prevalence = actual_{yes}/total$$

**Cross Validation** In k-fold cross-validation, the original sample is randomly partitioned into k equal sized subsamples. Of the k subsamples, a single subsample is retained as the validation data for testing the model, and the remaining k-1 subsamples are used as training data. The cross-validation process is then repeated k times, with each of the k subsamples used exactly once as the validation data. The k results can then be averaged to produce a single estimation. The advantage of this method over repeated random sub-sampling (see below) is that all observations are used for both training and validation, and each observation is used for validation exactly once.

```

1 from sklearn.model_selection import cross_validate
2 from sklearn import metrics
3
4 def scoringModel(model, train, y):
5     scores = cross_validate(model, train, y, scoring=scoring,
6                             cv=5)
7     return scores
8
9 scoreTreeModel = scoringModel(decisionTreeModel, train_X,
10                               train_y)
11 scoreLR = scoringModel(logisticRegressionModel, train_X,
12                         train_y)

```

```

11 data = {'Model' : ['Decision tree', 'Logistic regression']}
12
13 for metric_name in scoreTreeModel.keys():
14     data[metric_name] = [scoreTreeModel[metric_name].mean(),
15                          scoreLR[metric_name].mean()]
16
17 df = pd.DataFrame(data)
18 df.T

```

**Listing 1.1.** Scoring model function code

Because of low amount of data, we used Cross validation [4].

**Table 4.** Evaluation of models.

| Model                | Decision tree | Logistic regression |
|----------------------|---------------|---------------------|
| fit_time             | 0.00423088    | 0.00406737          |
| score_time           | 0.0128297     | 0.0135069           |
| test_accuracy        | 0.845275      | 0.790989            |
| test_precision_macro | 0.660598      | 0.407949            |
| test_recall_macro    | 0.665909      | 0.482576            |
| test_f1_weighted     | 0.82632       | 0.724681            |
| test_roc_auc         | 0.665909      | 0.291919            |

### 3 Improvement

#### 3.1 Normalization

Standardize features by removing the mean and scaling to unit variance

The standard score of a sample  $x$  is calculated as:

$$z = (x - u)/s$$

where  $u$  is the mean of the training samples or zero if  $with_{mean} = False$ , and  $s$  is the standard deviation of the training samples or one if  $with_{std} = False$ .

After normalization:

After normalization, fitting and scoring time increased and other performances are increased by little.

#### 3.2 Feature reduction

There are lot of techniques to reduce features. But we mentioned before, High correlated features are selected now along the correlation diagram such as age, time, type and induration<sub>diameter</sub>.

**Table 5.** Evaluation of models after Normalization

| Model                | Decision tree | Logistic regression |
|----------------------|---------------|---------------------|
| fit_time             | 0.00378752    | 0.0033298           |
| score_time           | 0.0120449     | 0.0117829           |
| test_accuracy        | 0.845275      | 0.776703            |
| test_precision_macro | 0.660598      | 0.406667            |
| test_recall_macro    | 0.665909      | 0.474242            |
| test_f1_weighted     | 0.82632       | 0.716681            |
| test_roc_auc         | 0.665909      | 0.502273            |

**Table 6.** Evaluation of models after Normalization and feature reduction

| Model                | Decision tree | Logistic regression |
|----------------------|---------------|---------------------|
| fit_time             | 0.00363851    | 0.003368            |
| score_time           | 0.0122293     | 0.0125299           |
| test_accuracy        | 0.860659      | 0.805275            |
| test_precision_macro | 0.66188       | 0.409048            |
| test_recall_macro    | 0.675         | 0.490909            |
| test_f1_weighted     | 0.834291      | 0.732066            |
| test_roc_auc         | 0.675         | 0.622222            |

**Conclusion and Future work** Even though we have too less number of samples for this work, We analyzed data and learned classification task. Moreover, Data over-sampling, finding outliers, feature engineering and hyper-parameter tuning can be helped to improve accuracy and performance. Most importantly, we should gather more samples for this case.

## References

1. UCI Immunotherapy data set,  
<https://archive.ics.uci.edu/ml/datasets/Immunotherapy+Dataset>. Last accessed 4 Jan 2020
2. Classification Algorithms in Machine Learning,  
<https://medium.com/datadriveninvestor/classification-algorithms-in-machine-learning-85c0ab65ff4>. Last accessed 4 Jan 2020
3. Logistic regression,  
[https://ml-cheatsheet.readthedocs.io/en/latest/logistic\\_regression.html](https://ml-cheatsheet.readthedocs.io/en/latest/logistic_regression.html). Last accessed 4 Jan 2020
4. Cross Validation  
[https://en.wikipedia.org/wiki/Cross-validation\\_\(statistics\)](https://en.wikipedia.org/wiki/Cross-validation_(statistics)). Last accessed 4 Jan 2020
5. Building A Logistic Regression in Python,  
<https://towardsdatascience.com/building-a-logistic-regression-in-python-step-by-step-becd4d56c9c8>. Last accessed 4 Jan 2020

## Appendix A

```

1 from sklearn import tree
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn.metrics import mean_absolute_error
4 import graphviz
5 from IPython.display import Image
6
7 decisionTreeModel = DecisionTreeClassifier(random_state=1)
8 decisionTreeModel.fit(train_X, train_y)
9
10 dot_data = tree.export_graphviz(decisionTreeModel, out_file=
    None,
11     feature_names = train_X.columns.values.tolist(),
12     class_names= ['0','1'],
13     filled=True,
14     rounded=True
15 )
16 graph = graphviz.Source(dot_data)
17 from IPython.display import SVG
18 Image(graph.pipe(format='png'))

```

Listing 1.2. Decision tree model visualization

```

1 from sklearn.model_selection import train_test_split
2
3 X = data.loc[:, data.columns != 'Result_of_Treatment']
4 y = data.loc[:, 'Result_of_Treatment']
5 train_X, test_X, train_y, test_y = train_test_split(X, y,
    test_size=0.2, random_state=1)
6 #train_X, val_X, train_y, val_y = train_test_split(train_X,
    train_y, test_size=0.2, random_state=1)
7
8 from sklearn.linear_model import LogisticRegression
9
10 LR = LogisticRegression()
11 scoreLR = scoring(LR)
12
13 data = {'Model' : ['Decision tree', 'Logistic regression']}
14
15 for metric_name in scoreTreeModel.keys():
16     data[metric_name] = [scoreTreeModel[metric_name].mean(),
        scoreLR[metric_name].mean()]
17
18 df = pd.DataFrame(data)
19 df.T

```

Listing 1.3. Training Logistic regression and Decision tree model

```

1
2 #check null values
3 cols_with_missing = [col for col in data.columns
4                       if data[col].isnull().any()]
5
6 pd.plotting.register_matplotlib_converters()
7 import matplotlib.pyplot as plt
8 %matplotlib inline
9 import seaborn as sns
10
11 sns.countplot(data['Result_of_Treatment'], label="Sum")
12 data['Result_of_Treatment'].value_counts()
13 #TODO We can use SMOTE algorithm for oversampling
14
15 pd.concat([data.groupby('Result_of_Treatment').mean(), data.
16            groupby('Result_of_Treatment').mean().diff().dropna()]).T
17
18 ax = sns.distplot( data[data['Result_of_Treatment'] == 0]['
19                    age'], color="skyblue", label="0")
20
21 ax = sns.distplot( data[data['Result_of_Treatment'] == 1]['
22                    age'], color="red", label="1")
23 ax.legend()
24
25 ax = sns.distplot( data[data['Result_of_Treatment'] == 0]['
26                    Time'], color="skyblue", label="0")
27 ax = sns.distplot( data[data['Result_of_Treatment'] == 1]['
28                    Time'], color="red", label="1")
29 ax.legend()
30
31 sns.catplot(x="Result_of_Treatment", y="Time", kind = 'violin
32             ', palette="pastel", inner="stick", data=data);
33 sns.catplot(x="Result_of_Treatment", y="age", kind = 'violin'
34             , palette="pastel", inner="stick", data=data);
35 sns.catplot(x="Type", y="Time", hue= 'Result_of_Treatment',
36             kind = 'violin', palette="pastel", inner="stick", split=
37             True, data=data);
38
39 corr = data.corr()
40 corr.T
41
42 mask = np.zeros_like(corr)
43 mask[np.triu_indices_from(mask)] = True
44 with sns.axes_style("white"):
45     f, ax = plt.subplots(figsize=(7, 5))
46     ax = sns.heatmap(corr, mask=mask, annot = True,
47                     vmin =-1, vmax =1, center = 0,
48                     xticklabels=corr.columns, cmap="RdBu_r")
49
50 sns.scatterplot(x=data['Time'], y=data['age'], hue=data['
51                 Result_of_Treatment'])

```



```

41
42 # 2D KDE plot
43 sns.jointplot(x=data['Time'], y=data['age'], kind="kde")
44
45 fig = plt.figure()
46 fig.subplots_adjust(hspace=1, wspace=.4)
47 ax = fig.add_subplot(1, 2, 1)
48 ax = sns.scatterplot(x=data['Time'], y=data['age'], hue=data[
    'Result_of_Treatment'], ax = ax)
49 ax = fig.add_subplot(1, 2, 2)
50 succeeded = data.loc[data.Result_of_Treatment == 1]
51 failed = data.loc[data.Result_of_Treatment == 0]
52 ax = sns.kdeplot(failed.Time, failed.age, label = "0", cmap="
    Blues", ax=ax)
53 ax = sns.kdeplot(succeeded.Time, succeeded.age, label = "1",
    cmap="Reds", ax = ax)
54 ax.legend()
55 plt.show()
56
57 succeeded = data.loc[data.Result_of_Treatment == 1]
58 failed = data.loc[data.Result_of_Treatment == 0]
59 ax = sns.kdeplot(failed.Time, failed.age, label = "0",
60                 cmap="Blues")
61 ax = sns.kdeplot(succeeded.Time, succeeded.age, label = "1",
62                 cmap="Reds")
63 ax.legend()
64
65 c = sns.FacetGrid(data, col="Type", hue="Result_of_Treatment"
66                  )
67 c.map(plt.scatter, "Time", "age", alpha=.7)
68 c.add_legend();

```

Listing 1.4. Data exploration and visualization code

```

1 from sklearn import preprocessing
2
3 std_scale = preprocessing.StandardScaler().fit(train_X)
4 x_train_norm = std_scale.transform(train_X)
5 training_norm_X = pd.DataFrame(x_train_norm, index=train_X.
6                               index, columns=train_X.columns)
7 print(training_norm_X.head())
8
9 x_test_norm = std_scale.transform(test_X)
10 testing_norm_X = pd.DataFrame(x_test_norm, index=test_X.index
11                               , columns=test_X.columns)
12 print(testing_norm_X.head())

```

Listing 1.5. Normalization

```

1 def sigmoid(z):
2     return 1.0 / (1 + np.exp(-z))
3
4 def predict(features, weights):
5     '''
6     Returns 1D array of probabilities
7     that the class label == 1
8     '''
9     z = np.dot(features, weights)
10    return sigmoid(z)
11
12 def cost_function(features, labels, weights):
13    '''
14    Using Mean Absolute Error
15    Features:(100,3)
16    Labels: (100,1)
17    Weights:(3,1)
18    Returns 1D matrix of predictions
19    Cost = (labels*log(predictions) + (1-labels)*log(1-
20    predictions) ) / len(labels)
21    '''
22    observations = len(labels)
23    predictions = predict(features, weights)
24    #Take the error when label=1
25    class1_cost = -labels*np.log(predictions)
26    #Take the error when label=0
27    class2_cost = (1-labels)*np.log(1-predictions)
28    #Take the sum of both costs
29    cost = class1_cost - class2_cost
30    #Take the average cost
31    cost = cost.sum() / observations
32
33    return cost
34
35 def update_weights(features, labels, weights, lr):
36    '''
37    Vectorized Gradient Descent
38
39    Features:(200, 3)
40    Labels: (200, 1)
41    Weights:(3, 1)
42    '''
43    N = len(features)
44    #1 - Get Predictions
45    predictions = predict(features, weights)
46    #2 Transpose features from (200, 3) to (3, 200)
47    gradient = np.dot(features.T, predictions - labels)
48    #3 Take the average cost derivative for each feature
49    gradient /= N
50    #4 - Multiply the gradient by our learning rate

```

```

50     gradient *= lr
51     #5 - Subtract from our weights to minimize cost
52     weights -= gradient
53
54     return weights
55 def decision_boundary(prob):
56     return 1 if prob >= .5 else 0
57
58 def classify(predictions):
59     '''
60     input - N element array of predictions between 0 and 1
61     output - N element array of 0s (False) and 1s (True)
62     '''
63     _decision_boundary = np.vectorize(decision_boundary)
64     return _decision_boundary(predictions).flatten()
65
66 def train(features, labels, weights, lr, iters):
67     cost_history = []
68
69     for i in range(iters):
70         #print (features)
71         #print (weights)
72
73         weights = update_weights(features, labels, weights,
74                                 lr)
75
76         #Calculate error for auditing purposes
77         cost = cost_function(features, labels, weights)
78         cost_history.append(cost)
79
80         # Log Progress
81         if i % 100 == 0:
82             print ("iter: {} cost: {}".format(i, cost))
83
84     return weights, cost_history
85
86 #CALCULATION
87 np.random.seed(0)
88 weights, cost_history = train(np.hstack((np.ones((train_y.
89 count(),1)), train_X[['Time', 'age']].to_numpy()))),
90                             train_y.to_numpy(),
91                             np.random.uniform(low=-1.0, high=1.0, size=3),
92                             0.001, 1000)
93 possibilities = predict(np.hstack((np.ones((test_y.count(),1))
94 ,test_X[['Time', 'age']].to_numpy()))), weights)
95 predictions = classify(possibilities)
96 possibilities

```

Listing 1.6. Logistic regression from scratch